

Feature Extraction

Feature Extraction for CNNs (Computer Vision)

In CNNs, feature extraction is the **primary function of the convolutional layers**. It is the process of automatically learning and creating **new, compact, and abstract representations** (features) from raw pixel data.

How it Works

1. **Convolutional Layers:** These layers use **learnable filters (kernels)** to scan the input image. Each filter is designed to detect a specific spatial pattern, such as a **vertical edge**, a **texture**, or a **corner**. The output of a convolutional layer, called a **feature map**, is a new, engineered feature that shows where that specific pattern exists in the image.
2. **Hierarchical Learning:**
 - o **Early Layers** extract simple, **generic features** (e.g., edges, curves).
 - o **Deeper Layers** combine these simple features to extract complex, **high-level semantic features** (e.g., an eye, a wheel, or a whole object).
3. **Transfer Learning:** In practical use, feature extraction often involves taking a CNN pre-trained on a massive dataset (like ImageNet), removing its final classification layer, and using the output of the last convolutional/pooling layer as the **extracted feature vector** for a new, related task. This vector is a highly condensed and useful set of features representing the original image.

Feature Extraction for LLMs (Natural Language Processing)

For Large Language Models based on the Transformer architecture, feature extraction is the process of converting raw text into **rich, context-aware numerical representations (vectors)** that capture the text's meaning, syntax, and semantics.

How it Works

1. **Tokenization and Embedding:** The raw text is first broken into tokens (words/sub-words). Each token is then mapped to an initial numerical vector called an **embedding**. This vector is a basic feature representing the word's semantic meaning.
2. **Contextualization (The Transformer Layers):** The subsequent stacked Transformer layers, powered by the **attention mechanism**, iteratively refine these initial features. At each layer, the model performs a calculation that effectively **creates a new feature vector** for every token by incorporating information from all other tokens in the sequence.
3. **Abstract Feature Vectors:** The final output of the model's layers for a given input text is

a sequence of **contextualized embeddings**. These are highly complex feature vectors where:

- The vector for the word "bank" has features that make it distinct in "river bank" versus "money bank."
- The vectors encode features like **part-of-speech**, **coreference** (which words refer to the same entity), and overall **sentiment** or topic.

This process transforms the sequence of raw tokens into a compact, fixed-size matrix of **linguistic features** that can be used directly for downstream tasks like classification or sequence generation.

Vs. Feature Selection

Feature Selection for CNNs (Computer Vision)

In the context of CNNs, feature selection refers to choosing which learned **visual features** are most effective for a specific task. However, the primary challenge isn't typically *selecting* features but rather determining the optimal **architecture** or **transfer learning strategy** to extract the right features.

How it Works

A CNN is structured to automatically extract features from raw pixel data across its layers:

1. **Early Layers (Generic Features):** These layers learn **low-level, generic features** like edges, corners, and color blobs. These features are useful for almost any image task.
2. **Middle/Later Layers (Specific Features):** As the network deepens, the receptive fields grow, and the layers learn **high-level, semantic features** specific to the training data (e.g., eyes and ears in a face detection network).

Selection Strategies (Transfer Learning)

Feature selection in CNNs is often implemented through **Transfer Learning**:

- **Feature Extraction (Freezing):** The most common selection method is to use a pre-trained CNN (like VGG or ResNet) and **freeze** (lock) all its weights up to a certain point. The output of the last frozen layer is taken as the set of **selected features** (e.g., a 2048-dimensional vector). These features are then fed into a new, smaller classifier for the specific task.
- **Layer Selection:** You are essentially *selecting* which set of learned features to use by choosing which layer's output to take. For generic tasks, you might select features from a lower layer; for specialized tasks, you select features from a higher layer that has learned more task-specific concepts

Feature Selection for LLMs (Natural Language Processing)

For Large Language Models (LLMs), feature selection takes place at two levels: the **input (text)** level and the **internal (model)** level. The most important selection process is handled internally by the **Attention Mechanism**.

How it Works

LLMs process text by converting it into numerical feature vectors (**embeddings**), which are then refined through stacked Transformer layers. The goal of feature selection here is to ensure the model focuses only on the most predictive linguistic signals.

1. **Input Feature Selection:** This is the traditional feature selection you might use before the model, such as preprocessing the text to remove stop words or using a statistical test to select the most relevant *tokens* (words) to include in the prompt.
2. **Internal Feature Selection (Attention):** This is the most crucial part. The self-attention mechanism performs dynamic feature selection for every *token* at every layer:
 - o **Dynamic Weighting:** When the model processes a token, the attention mechanism **selectively weighs** all other tokens in the context (the input or the text generated so far).
 - o **Context Filtering:** If a token is irrelevant to the meaning of the current word (e.g., if the word "river" is processing, the word "blue" in the context gets a high weight, while "yesterday" gets a low weight), the attention weight for the irrelevant token is effectively zeroed out. This is an extremely sophisticated form of feature selection, as the model **selects which parts of the context** (which features) to use dynamically, based on the task and the query.

The final output of the LLM's layers are **contextualized feature vectors** that encode rich information, having already had irrelevant information filtered out by the attention mechanism.